

Defeating SkyNet: Design Document

Stephen Tridgell, Nicholas Carydis

Botnets have had a prominence in the media over recent years, effectively instilling them into the public consciousness as another threat to the security of computer networks. The security protocols of an established botnet, although rarely addressed by the media, are an interesting and informative way of looking at the threat and of determining ways of having it neutralised.

The first part of the project involved securing an established botnet, by developing protocols to provide authentication, confidentiality, integrity and replay prevention between infected hosts. Design choices made and the reasons for them are outlined below.

- **What was your choice of Diffie-Hellman key exchange parameters and what made you select them specifically?**

The parameters chosen were those listed as part of the 2048 MODP group of the IETF standard (RFC 3526). The default allocated prime was increased in bit size to 2048 bits, with a hexadecimal value of

```
“      FFFFFFFF FFFFFFFF C90FDAA2 2168C234 C4C6628B 80DC1CD1
      29024E08 8A67CC74 020BBEA6 3B139B22 514A0879 8E3404DD
      EF9519B3 CD3A431B 302B0A6D F25F1437 4FE1356D 6D51C245
      E485B576 625E7EC6 F44C42E9 A637ED6B 0BFF5CB6 F406B7ED
      EE386BFB 5A899FA5 AE9F2411 7C4B1FE6 49286651 ECE45B3D
      C2007CB8 A163BF05 98DA4836 1C55D39A 69163FA8 FD24CF5F
      83655D23 DCA3AD96 1C62F356 208552BB 9ED52907 7096966D
      670C354E 4ABC9804 F1746C08 CA18217C 32905E46 2E36CE3B
      E39E772C 180E8603 9B2783A2 EC07A28F B5C55DF0 6F4C52C9
      DE2BCBF6 95581718 3995497C EA956AE5 15D22618 98FA0510
      15728E5A 8AACAA68 FFFFFFFF FFFFFFFF
“
```

The generator (or ‘g’) was selected to be 2, and the PyCrypto secure Random function was used to generate a random ‘a’ and ‘b’ for the two parties. The larger prime was selected to make a brute-force attack less feasible via the avenue of the discrete-logarithm problem, for a negligible cost of resources. This did not, however, increase the strength of the resultant symmetric key, as the shared secret is hashed using SHA256 to produce an un-biased 256b symmetric key.

- **What was your choice of cipher? What mode of operation does it use?**

The AES cipher was chosen with a symmetric key size of 256 bits. This key was produced by hashing the shared secret produced by the Diffie-Hellman key exchange. The choice of key size was based on the largest possible key supported by AES, so as to provide the maximum security possible at this juncture. The cipher was chosen to run in Cipher-Feedback mode (or CFB) so as to perform as a stream cipher, for the issuance of commands and reception of data between infected hosts.

For authentication purposes, a public-key cryptosystem (RSA) was used. PKCS1_OAEP (Optimal Asymmetric Encryption Padding) was used for the production of a public/private key pair, which was the standard RSA algorithm combined with a padding scheme to mitigate the risks of frequency

cryptanalysis on the resultant cyphertext.

- **How do you prevent attackers from tampering with messages in transit?**

The authentication procedure between client and server involves the use of an RSA exchange. When the commander sends a request to a host, the host replies with a challenge; This challenge includes a random IV along with one half of the DH secret (g^a). This challenge is encrypted using the parent's public key (known by the host), providing confidentiality and integrity. The controller replies with its other half of the secret along with the IV, which it signs with its private key – providing integrity (but not confidentiality in this case).

For other commands and data, the data are appended with its hash, which is then encrypted using AES. This provides confidentiality and integrity.

- **How do you prevent replay attacks?**

In the authentication stage, the random IV provides protection against replay attacks, as it changes every time the authentication protocol is invoked. For other communications, the use of CFB prevents replay attacks as a) the CFB system begins with the IV (which changes with each new connection) and b) future cyphertext outputs depend on the encrypted value of the previous blocks of cyphertext (for which an attacker would need to use the key to predict). Thus sending the same block of cyphertext twice will not decrypt to the same plaintext.

- **Why might we want to allow for peer-to-peer file transfers between bots? Etc.**

Using p2p file transfers allows the bots to run in a distributed fashion, without a need for a central controller. The advantages of using a centralised system are generally the use of less complex authentication mechanisms, as well as more judicious control over the bots in the network. Instructions and data would not require propagation through a hierarchy of hosts, and so would reach their destination faster too.

However, the disadvantages are numerous: It produces a single point of failure, and would pose an easy target for governments and other hackers to take down or commandeer the botnet. It would be much easier to trace commands emanating from a single web server, and for that web server to be taken down – either by redirecting its traffic at lower levels or DDOSing it. The server itself may not be able to cope with capacity demands as more and more hosts are infected and become part of the botnet – it may end up DOSing itself.

- **Explain how this botnet, if used in the real world, could be trivially controlled by other hackers and government agencies. How might one attempt to stop it?**

As mentioned, this centralised botnet could be trivially controlled once the central web server is compromised. The modifications we have made (including decentralising the structure into a tree-like hierarchy) defend against some of these attacks and make the root node harder to trace. However, even in this structure, a compromised node will effectively compromise those nodes beneath it, and compromising the root node will compromise the entire botnet.